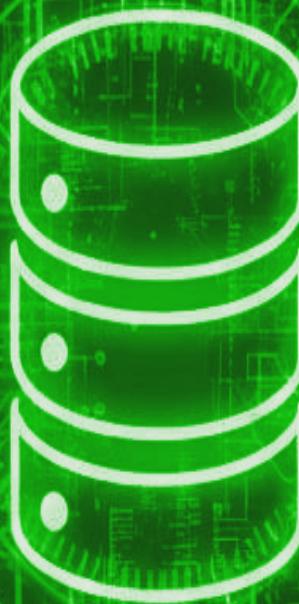


1. SUB-QURY
2. MAX & MIN
3. SUB-QURY TYPES
4. SUB-QURY OPERATORS
5. NETSDURY SUB-QUIRY
6. JOINS
7. TYPES OF JOINS
8. CORRLATED SUB-QURY
9. SUB-QURY VS CORRORIATED
10. EXISTS & NON-EXISTS
11. SINGLE ROW FUNCTIONS
13. NORMALIZATION
14. NORMAL FORM
15. NORMAL FORM LEVELS



## 1. Sub-Query

**Definition:** A subquery (or inner query) is a SQL query nested inside another query (outer query). It's enclosed in parentheses and returns data used by the outer query.

**Uses:** Filtering results (e.g., in WHERE clause), computing aggregates for comparisons, or deriving columns in SELECT. Professionals use them for complex data analysis, like finding outliers or conditional aggregates.

### Working Procedure:

1. The inner query executes first, producing a result set (single value, list, or table).
2. The outer query uses this result to complete its execution.
3. If the subquery depends on the outer, it's correlated (covered later).

**Causing Effects:** Can slow queries if not optimized (e.g., repeated execution), but improves readability. Poor design leads to high CPU usage; indexing helps mitigate.

### Examples:

Assume Employees table:

	emp_id	name	salary	dept_id
	1	Alice	50000	101
	2	Bob	60000	102
	3	Charlie	55000	101

Example: Find employees with salary above the average.

```
SELECT name, salary  
FROM Employees  
WHERE salary > (SELECT AVG(salary) FROM Employees);
```

Inner query calculates  $\text{AVG}(\text{salary}) = 55000$ . Outer filters names/salaries  $> 55000$  Raja

## 2. MAX & MIN Functions

**Definition:** Aggregate functions that return the highest (MAX) or lowest (MIN) value in a column or expression.

**Uses:** Summarizing data, like finding top earners or oldest records. Pros use with GROUP BY for analytics dashboards.

### Working Procedure:

1. Scan the column/expression.
2. Ignore NULLs unless specified.
3. Return single max/min value per group if grouped.

### Examples:

Using Employees table above.

Example 1 : Highest salary.

```
SELECT MAX(salary) AS highest_salary FROM Employees;
```

Scans salary column, picks max.

Example 2: Min salary per department.

```
SELECT dept_id, MIN(salary) AS min_sal  
FROM Employees  
GROUP BY dept_id;
```

Explanation: Groups by dept\_id, finds min per group.

```

Select
    CustomerId, Year(OrderDate) [Year], Month(OrderDate) [Month],
    Min(TotalAmount) as Monthly_Minimum,
    Max(TotalAmount) as Monthly_Maximum
From SalesOrders
Group By
    CustomerId, Year(OrderDate), Month(OrderDate)

```

100 %

	CustomerId	Year	Month	Monthly_Minimum	Monthly_Maximum
1	101	2018	1	75.000	220.000
2	101	2018	2	235.000	235.000
3	250	2018	1	100.000	230.000
4	250	2018	2	245.000	245.000
5	300	2018	1	140.000	275.000
6	300	2018	2	225.000	225.000

### 3. Types of Sub-Query

**Definition:** Subqueries classified by return type: Single-row (one value), Multiple-row (list), Multiple-column (table-like), Scalar (single value, like aggregates).

**Uses:** Single-row for comparisons; multiple-row with IN/ANY; scalar in SELECT for derived fields.

**Working Procedure:** Based on type, outer query operators adapt (e.g., = for single, IN for multiple).

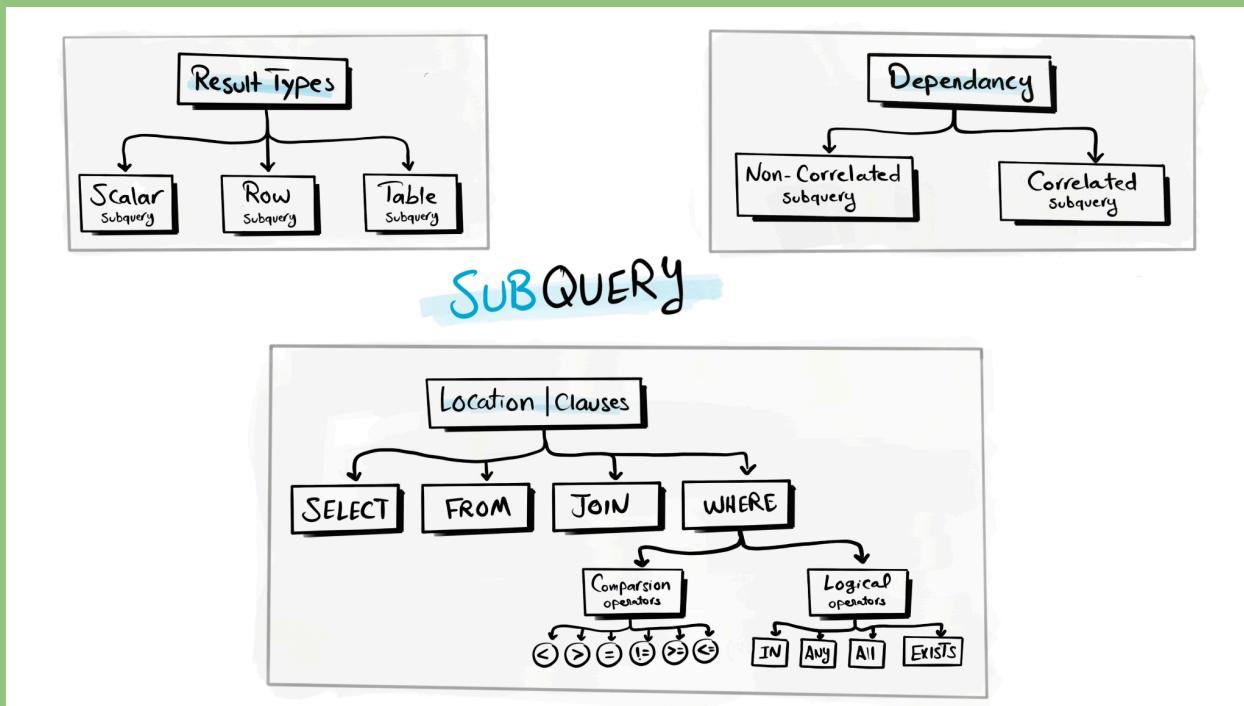
#### Examples:

Single-row: `SELECT AVG(salary) FROM Employees;`

Multiple-row:

```
SELECT name FROM Employees WHERE dept_id IN (SELECT dept_id FROM Departments
WHERE dept_name = 'Sales');
```

Explanation: Inner returns a list of dept\_ids; outer filters names.



#### 4. Sub-Query Operators

**Definition:** Operators like =, >, IN, ANY, ALL, and EXISTS are used with subqueries for comparisons.

**Uses:** IN for lists; ANY/ALL for conditional extrema; EXISTS for presence checks.

**Working Procedure:** Operator evaluates outer row against subquery result (e.g., IN checks membership).

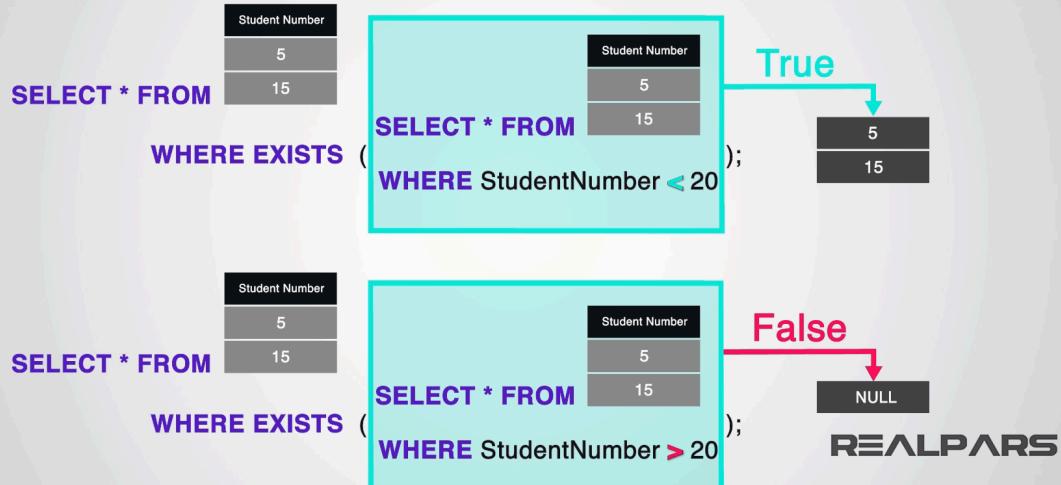
**Examples:**

ANY example:

```
SELECT name FROM Employees WHERE salary > ANY (SELECT salary FROM Employees WHERE dept_id = 101);
```

Explanation: > any value from subquery (e.g., >50000 or >55000), so Bob and Charlie if salaries vary.

## EXIST Operation



## 5. Nested Sub-Query

**Definition:** A subquery inside another subquery, creating multiple levels of nesting.

**Uses:** Complex filtering, like multi-level aggregations (e.g., avg of maxes).

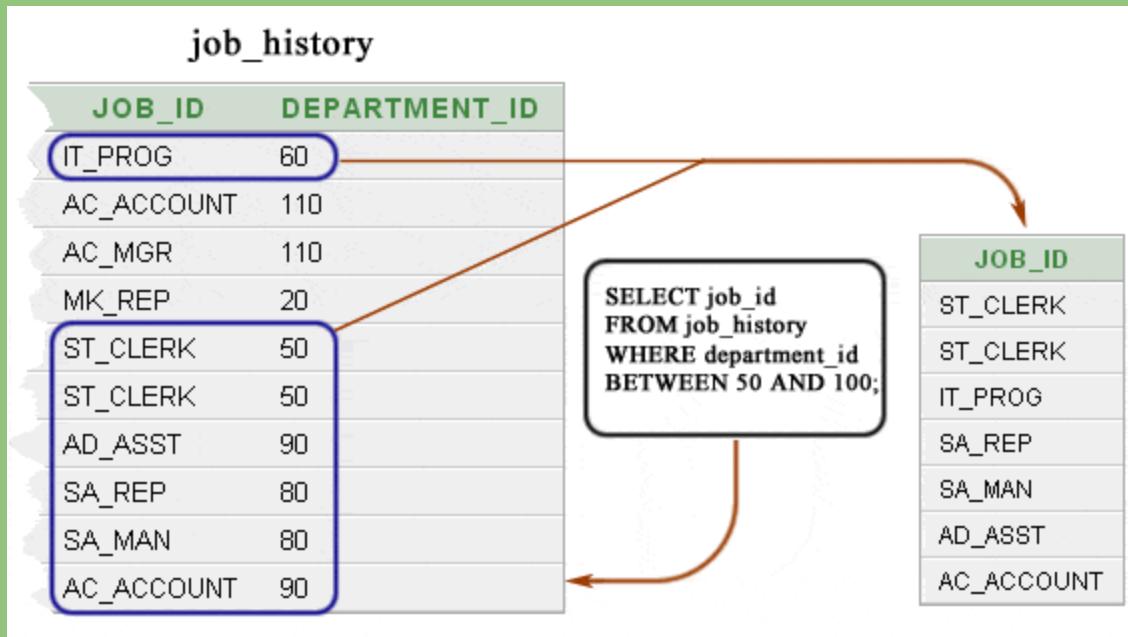
**Working Procedure:**

1. Innermost executes first.
2. Results propagate outward.

**Examples:**

```
SELECT name FROM Employees WHERE salary > (SELECT AVG(salary) FROM Employees
WHERE dept_id = (SELECT dept_id FROM Departments WHERE dept_name = 'IT'));
```

Explanation: Innermost finds IT dept\_id; middle avgs salaries there; outer filters.



## 6. Joins

**Definition:** Clauses to combine rows from two or more tables based on related columns.

**Uses:** Merging data across tables for reports and analysis.

**Working Procedure:**

1. Specify tables and join condition (ON).
2. Match rows; output combined.

**Examples:**

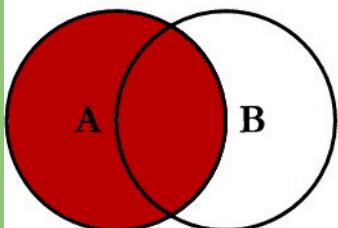
Basic INNER:

```
SELECT e.name, d.dept_name
```

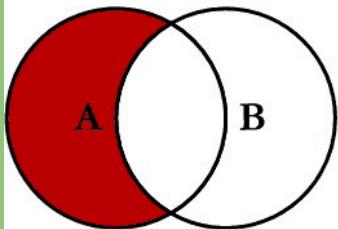
```
FROM Employees e JOIN Departments d ON e.dept_id = d.dept_id;
```

Explanation: Matches dept\_id, outputs names with depts.

# SQL JOINS

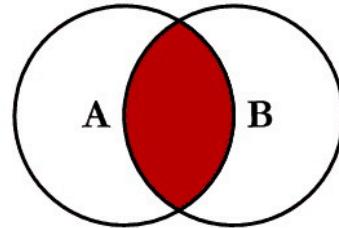


```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```

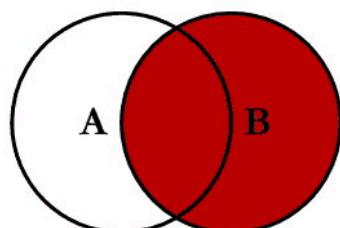


```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```

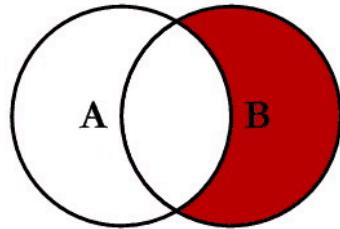
```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```

```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

© C.L. Moffatt, 2008

## 7. Types of Joins

**Definition:** INNER (matching rows), LEFT (all left + matching right), RIGHT (all right + matching left), FULL (all from both), CROSS (Cartesian), SELF (same table).

**Uses:** INNER for intersections; OUTER for inclusives; CROSS for combinations.

**Working Procedure:** Based on type, include unmatched rows with NULLs (OUTER) or all combos (CROSS).

**Differences with Others:** INNER vs. OUTER: Matching only vs. inclusive. SELF vs. others: Intra-table.

**Causing Effects:** OUTER can bloat results; CROSS explodes on large tables.

**Examples:**

LEFT JOIN:

```
SELECT e.name, d.dept_name
```

```
FROM Employees e LEFT JOIN Departments d ON e.dept_id = d.dept_id;
```

Explanation: All employees + dept names; NULL if no dept.

Type	Description	Matches Unmatched?
INNER	Only matching rows	No
LEFT	All left + matching right	Left yes, right no
RIGHT	All right + matching left	Right yes, left no
FULL	All from both	Yes
CROSS	All combinations	N/A
SELF	Join table to itself	As per condition

## 8. Correlated Sub-query

**Definition:** A subquery referencing columns from the outer query, executing per outer row.

**Uses:** Row-by-row comparisons, like finding relative maxima.

**Working Procedure:**

1. Outer query starts a row.
2. Inner references outer value, executes.
3. Repeat for each outer row.

**Causing Effects:** Slower ( $N \times M$  executions); good for small outer sets, bad for large (use joins instead).

**Examples:**

```
SELECT name, salary
```

```
FROM Employees e1
```

```
WHERE salary > (SELECT AVG(salary) FROM Employees e2 WHERE e1.dept_id =  
e2.dept_id);
```

## 9. Difference b/w Sub-query and Correlated Sub-query

Standard subquery independent; correlated depends on outer.

CTE VS CORRELATED SUBQUERY		
Aspect	CTE	Correlated Subquery
How It Runs	Runs once, and you can reuse the result.	Runs for each row in the outer query.
Readability	Super clear, especially for complex stuff.	A bit trickier to read since it's nested.
Speed	Great for big datasets if you reuse it.	Can be slower with lots of rows.
Reusability	Use it as many times as you want!	Runs fresh each time you call it.
Scope	Available for the whole query.	Only works where it's written.
Dependency	Stands on its own.	Needs the outer query to work.

## 10. Exists & Non-exists Operators

The EXISTS and NOT EXISTS operators are logical operators used in SQL to check for the presence or absence of records within a subquery. Unlike other operators that compare specific values, these operators evaluate to a boolean TRUE or FALSE based solely on whether the subquery returns any rows.

### 1. EXISTS Operator

The EXISTS operator returns TRUE if the subquery returns at least one row.

Purpose: Used to identify records in an outer query that have a corresponding relationship in another table.

Short-Circuiting: For efficiency, the database stops searching as soon as it finds the first matching row.

Example: Finding customers who have placed at least one order.

#### 1. EXISTS:

```
SELECT customer_name  
FROM Customers c  
WHERE EXISTS (SELECT 1 FROM Orders o WHERE o.customer_id = c.id);
```

### 2. NOT EXISTS Operator

The NOT EXISTS operator returns TRUE only if the subquery returns zero rows.

Purpose: Used to find "missing" data records in an outer query that have no related data in the subquery.

Example: Finding products that have never been sold.

```
SELECT product_name  
FROM Products p  
WHERE NOT EXISTS (SELECT 1 FROM Sales s WHERE s.product_id = p.id);
```

# EXISTS & NOT EXISTS in SQL

@lasha-dolenjashvili

```
SELECT * FROM Customers C
WHERE EXISTS (
    SELECT 1 FROM Payments P
    WHERE C.CustomerID = P.CustomerID
);
```

## Single Row Functions

Single-row functions (also known as scalar functions) are built-in SQL tools that operate on individual data items and return exactly one result for every row processed. Unlike aggregate functions, which summarize multiple rows into one result, single-row functions act on each row independently.

### Key Characteristics

One result per row: They return one output value for every row in the result set.

Versatility: They can be used in SELECT lists, WHERE clauses, ORDER BY clauses, and even HAVING clauses.

Nesting: These functions can be nested within one another to perform complex calculations in a single step.

Arguments: They can accept columns, literal values, or expressions as input arguments.

### Common Categories and Examples

## **1. Character Functions**

Used to manipulate text strings or change their case.

Case Conversion: UPPER('text') (TEXT), LOWER('TEXT') (text), and INITCAP('text') (Text).

Manipulation:

CONCAT(s1, s2): Joins two strings together.

SUBSTR(str, start, len): Extracts a portion of a string.

LENGTH(str): Returns the number of characters in a string.

TRIM(str): Removes leading or trailing spaces.

REPLACE(str, search, replace): Swaps specific text within a string.

Ex: Query: Display employee names in uppercase and calculate the length of their names.

SELECT

```
UPPER(first_name) AS loud_name,  
LENGTH(first_name) AS name_length  
FROM employees;
```

Result: If the name is "Alice", it returns ALICE and 5.

## **2. Numeric Functions**

Accept numeric input and return a single numeric value.

ROUND(n, d): Rounds a number to a specified number of decimal places.

TRUNC(n, d): Truncates a number without rounding.

MOD(n1, n2): Returns the remainder of a division operation.

CEIL(n) / FLOOR(n): Rounds up or down to the nearest whole integer.

Ex: Query: Round the unit price to two decimal places and calculate the remainder of an ID division.

```
SELECT
```

```
product_name,  
ROUND(unit_price, 2) AS price,  
MOD(product_id, 2) AS id_parity  
FROM products;
```

Result: If the price is 10.556, it returns 10.56.

### 3. Date Functions

Used to perform arithmetic or formatting on date and time values.

SYSDATE: Returns the current system date and time.

MONTHS\_BETWEEN(d1, d2): Calculates the count of months between two dates.

ADD\_MONTHS(d, n): Adds a specific number of months to a date.

LAST\_DAY(d): Returns the last day of the month for a given date.

Ex: Query: Show how many months an employee has been with the company as of today.

```
SELECT
```

```
last_name,  
MONTHS_BETWEEN(SYSDATE, hire_date) AS tenure_months  
FROM employees;
```

Result: Returns a numeric value representing the total months since the hire date.

## 4. Conversion Functions

Convert data from one data type to another.

TO\_CHAR(n/d, format): Converts a number or date into a formatted string.

TO\_NUMBER(str): Converts a numeric string into an actual number.

TO\_DATE(str, format): Converts a string into a date format the database recognizes.

Ex: Query: Convert a raw date into a specific readable string format.

SELECT

```
order_id,  
TO_CHAR(order_date, 'Day, DD Month YYYY') AS formatted_date  
FROM orders;
```

Result: Converts a date object to text like Saturday, 24 January 2026.

## 5. General & NULL Functions

Used primarily for handling NULL values and conditional logic.

NVL(val, default): Replaces a NULL with a specific value.

COALESCE(list): Returns the first non-null value from a list of arguments.

NULLIF(val1, val2): Returns NULL if two values are equal; otherwise returns the first value.

Ex: Query: Calculate total compensation by replacing NULL commissions with 0.

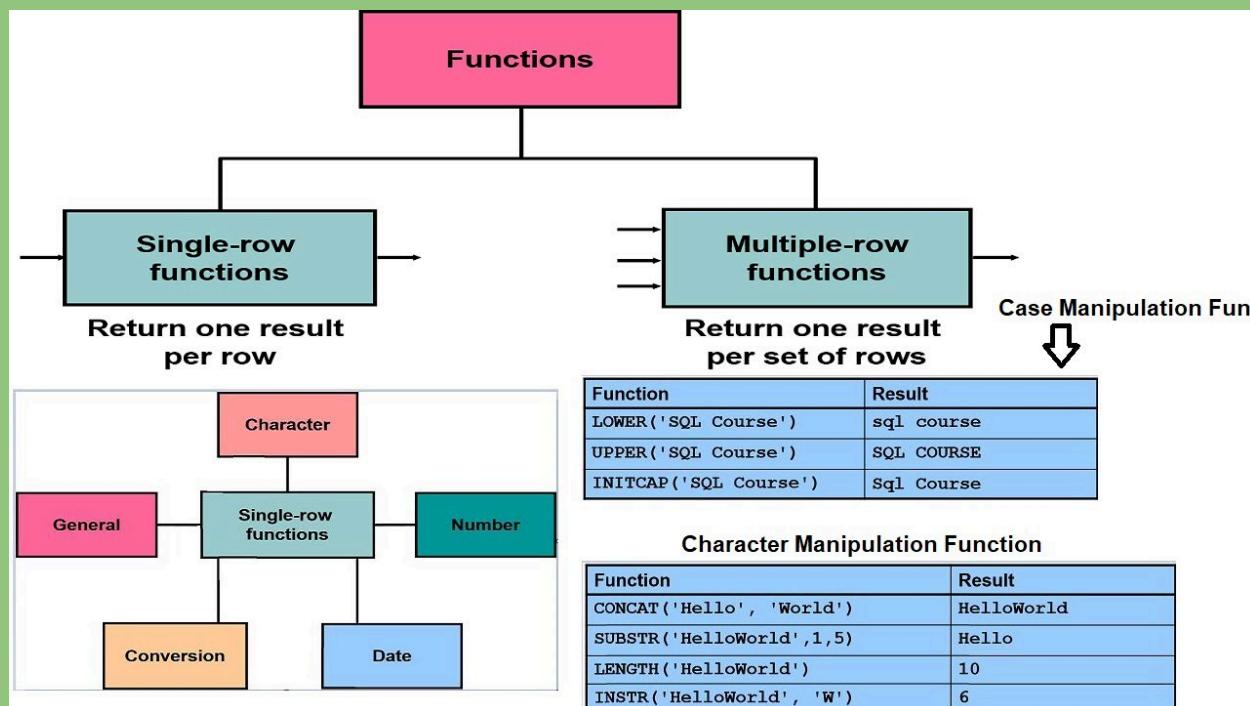
SELECT

```
last_name,  
salary + NVL(commission_pct, 0) AS total_pay  
FROM employees;
```

Result: If commission\_pct is NULL, the math still works because NVL treats it as 0.

Table of examples:

Function	Example SQL	Output (for Alice)
UPPER	UPPER(name)	ALICE
CONCAT	CONCAT(name, ' earns ', salary)	Alice earns 50000
ROUND	ROUND(salary * 1.1)	55000
NVL	NVL(null_column, 'Default')	Default



## Key Comparison

Feature	Single-Row Functions	Aggregate Functions
Input	One row at a time	A group of rows
Output	One result per row	One result per group
Example	UPPER, ROUND, TO_CHAR	SUM, AVG, COUNT
Usage	SELECT, WHERE, ORDER BY	SELECT, HAVING

## Statements in SQL

SQL statements are instructions used to communicate with and manipulate relational databases.

### Statements Classified into 5 Different Types (DDL, DML, TCL, DCL, DQL)

- **DDL (Data Definition Language):** Defines structures (CREATE, ALTER, DROP).
- **DML (Data Manipulation Language):** Manipulates data (INSERT, UPDATE, DELETE).
- **TCL (Transaction Control Language):** Manages transactions (COMMIT, ROLLBACK).
- **DCL (Data Control Language):** Controls access (GRANT, REVOKE).
- **DQL (Data Query Language):** Retrieves data (SELECT).

#### 1. Data Query Language (DQL)

Used exclusively for retrieving data from the database.

**SELECT:** The most fundamental statement, used to fetch specific data from one or more tables.

Example: SELECT name, age FROM users WHERE city = 'New York';

## **2. Data Manipulation Language (DML)**

DML is used to add, modify, or remove data records within existing table structures.

- **INSERT:** Adds new rows of data.
  - Example: `INSERT INTO students (id, name, age) VALUES (1, 'Rahul', 20);`
- **UPDATE:** Changes existing data.
  - Example: `UPDATE students SET age = 21 WHERE id = 1;`
- **DELETE:** Removes specific records based on a condition.
  - Example: `DELETE FROM students WHERE id = 1;`

**MERGE:** Performs "upsert" operations (updates existing rows or inserts new ones).

## **3. Data Definition Language (DDL)**

DDL defines and manages the physical structure (schema) of database objects like tables and indexes.

- **CREATE:** Builds a new object.
  - Example: `CREATE TABLE students (id INT PRIMARY KEY, name VARCHAR(50), age INT);`
- **ALTER:** Modifies an existing structure.
  - Example: `ALTER TABLE students ADD email VARCHAR(100);`
- **DROP:** Deletes an entire object and its data permanently.
  - Example: `DROP TABLE students;`
- **TRUNCATE:** Empties all records from a table while keeping its structure intact.
  - Example: `TRUNCATE TABLE students;`

## **4. Data Control Language (DCL)**

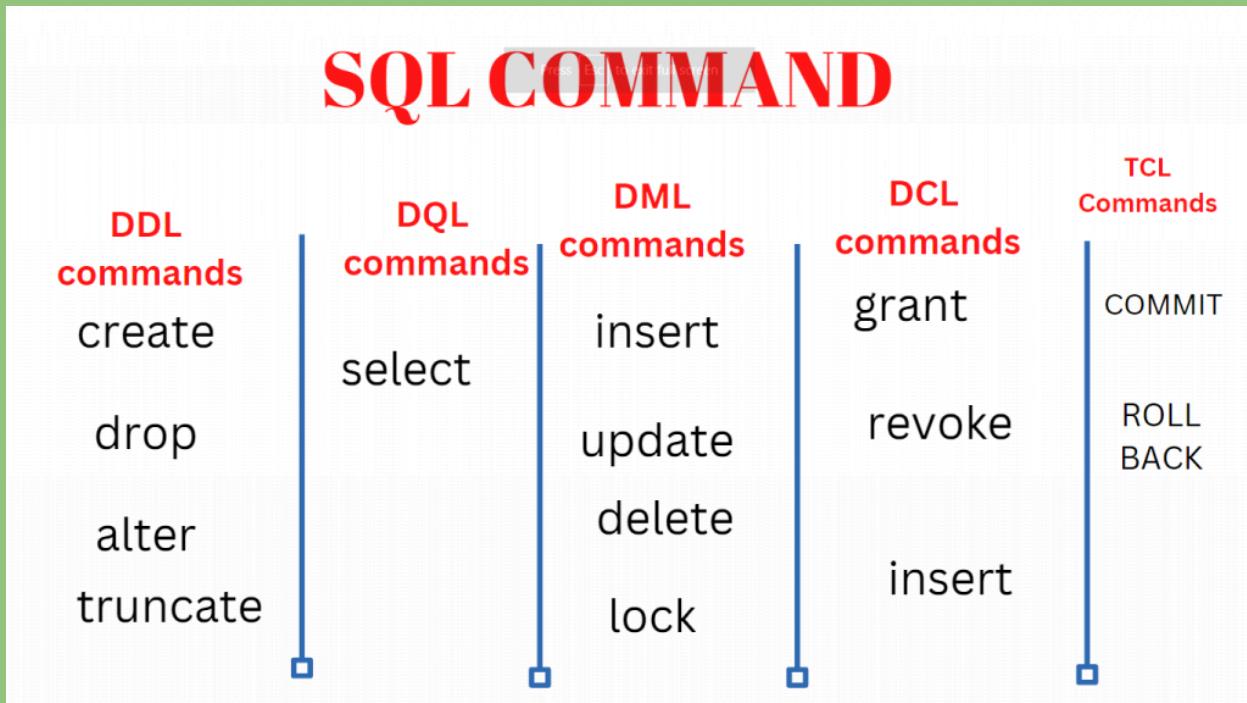
DCL manages security by controlling who can access or modify specific data.

- **GRANT:** Gives a user specific permissions.
  - Example: `GRANT SELECT, INSERT ON students TO user1;`
- **REVOKE:** Withdraws previously granted permissions.
  - Example: `REVOKE SELECT ON students FROM user1;`

## 5. Transaction Control Language (TCL)

TCL ensures data integrity by managing groups of DML operations as "transactions".

- COMMIT: Permanently saves all changes made during the current transaction.
  - Example: COMMIT;
- ROLLBACK: Reverts changes to the last committed state if an error occurs.
  - Example: ROLLBACK;
- SAVEPOINT: Sets a marker within a transaction to allow for partial rollbacks.
  - Example: SAVEPOINT point1;



### **13. Normalization**

" It is the process of reducing a large table into smaller tables in order to remove redundancies and anomalies by identifying their functional dependencies is known as Normalization. "

Or

"The process of decomposing a large table into smaller tables is known as Normalization ."

Or

"Reducing a table to its Normal Form is known as Normalization ."

**Why Normalization is Needed**

**Reduces Data Redundancy:** Prevents the same data from being stored in multiple places.

**Eliminates Anomalies:** Fixes issues where data cannot be inserted, updated, or deleted properly.

**Ensures Data Integrity:** Reduces inconsistencies, as updating a value only needs to be done in one place.

**Optimizes Structure:** Organizes data logically, often following the Single Responsibility Principle.

### **Normal Forms (Levels of Normalization)**

Normalization is performed in stages, known as Normal Forms:

**1NF (First Normal Form):** Ensures that all table columns contain atomic (indivisible) values and there are no repeating groups.

**2NF (Second Normal Form):** Must be in 1NF and ensure that all non-key attributes are fully functionally dependent on the primary key (no partial dependencies).

**3NF (Third Normal Form):** Must be in 2NF and ensure no transitive dependencies exist (non-key attributes should not depend on other non-key attributes).

**BCNF (Boyce-Codd Normal Form):** A stronger version of 3NF, often called 3.5NF, where for every functional dependency  $X \rightarrow Y$ , X is a super key.

Normal forms are a set of progressive rules (or design checkpoints) for relational schemas that reduce redundancy and prevent data anomalies. Each normal form - 1NF, 2NF, 3NF, BCNF, 4NF, 5NF - is stricter than the previous one: meeting a higher normal form implies the lower ones are satisfied. Think of them as layers of cleanliness for your tables: the deeper you go, the fewer redundancy and integrity problems you'll have.

### **Benefits of using Normal Forms**

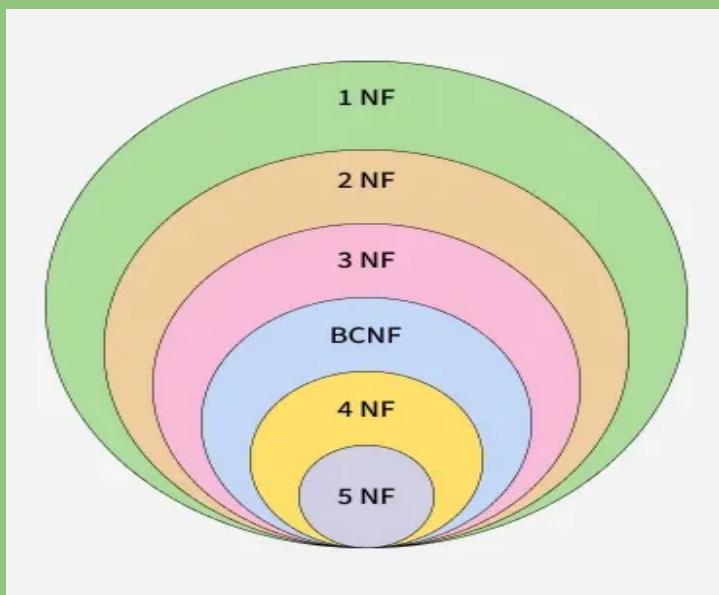
Reduce duplicate data and wasted storage.

Prevent insert, update, and delete anomalies.

Improve data consistency and integrity.

Make the schema easier to maintain and evolve.

The Diagram below shows the hierarchy of database normal forms. Each inner circle represents a stricter level of normalization, starting from 1NF (basic structure) to 5NF (most refined). As you move inward, data redundancy reduces and data integrity improves. Each level builds upon the previous one to ensure a cleaner and more efficient database design.



## **1. First Normal Form (1NF): Eliminating Duplicate Records**

A table is in 1NF if it satisfies the following conditions:

All columns contain atomic values (i.e., indivisible values).

Each row is unique (i.e., no duplicate rows).

Each column has a unique name.

The order in which data is stored does not matter.

Example of 1NF Violation: If a table has a column "Phone Numbers" that stores multiple phone numbers in a single cell, it violates 1NF. To bring it into 1NF, you need to separate phone numbers into individual rows.

- Problem: Storing multiple items in one cell (e.g., Items: 'Milk, Bread').
- SQL Example:

//Violates 1NF: Multiple values in one column

```
CREATE TABLE Orders (OrderID INT, Items VARCHAR(100));
```

//Corrected 1NF: Each item gets its own row

```
CREATE TABLE Orders_1NF (OrderID INT, Item VARCHAR(50));
```

```
INSERT INTO Orders_1NF VALUES (101, 'Milk'), (101, 'Bread');
```

## **2. Second Normal Form (2NF): Eliminating Partial Dependency**

A relation is in 2NF if it satisfies the conditions of 1NF and additionally. No partial dependency exists, meaning every non-prime attribute (non-key attribute) must depend on the entire primary key, not just a part of it.

Example: For a composite key (StudentID, CourseID), if the "StudentName" depends only on "StudentID" and not on the entire key, it violates 2NF. To normalize, move StudentName into a separate table where it depends only on "StudentID".

- Problem: In a table with a composite key (StudentID, CourseID), the CourseName depends only on CourseID.
- SQL Example:

//Corrected 2NF: Splitting tables so CourseName is in its own table

```
CREATE TABLE Courses (CourseID INT PRIMARY KEY, CourseName VARCHAR(50));
```

```
CREATE TABLE Student_Enrollment (StudentID INT, CourseID INT, PRIMARY KEY(StudentID, CourseID));
```

### 3. Third Normal Form (3NF): Eliminating Transitive Dependency

A relation is in 3NF if it satisfies 2NF and additionally, there are no transitive dependencies. In simpler terms, non-prime attributes should not depend on other non-prime attributes.

Example: Consider a table with (StudentID, CourseID, Instructor). If Instructor depends on "CourseID", and "CourseID" depends on "StudentID", then Instructor indirectly depends on "StudentID", which violates 3NF. To resolve this, place Instructor in a separate table linked by "CourseID".

- Problem: If ZipCode determines City, and City is in the Employees table, the City depends on the ZipCode rather than the EmployeeID.
- SQL Example:

//Corrected 3NF: Move Zip/City to a separate lookup table

```
CREATE TABLE Locations (ZipCode INT PRIMARY KEY, City VARCHAR(50));
```

```
CREATE TABLE Employees (EmpID INT PRIMARY KEY, Name VARCHAR(50), ZipCode INT);
```

### 4. Boyce-Codd Normal Form (BCNF): The Strongest Form of 3NF

BCNF is a stricter version of 3NF where for every non-trivial functional dependency ( $X \rightarrow Y$ ), X must be a superkey (a unique identifier for a record in the table).

Example: If a table has a dependency (StudentID, CourseID)  $\rightarrow$  Instructor, but neither "StudentID" nor "CourseID" is a superkey, then it violates BCNF. To bring it into BCNF, decompose the table so that each determinant is a candidate key.

- Problem: A table has multiple candidate keys that overlap.
- SQL Example:

```
//BCNF ensures that if a Professor determines a Subject,  
//the Professor table only maps ProfessorID to SubjectID.  
  
CREATE TABLE Professor_Subject (ProfessorID INT PRIMARY KEY, SubjectID INT);
```

## 5. Fourth Normal Form (4NF): Removing Multi-Valued Dependencies

A table is in 4NF if it is in BCNF and has no multi-valued dependencies. A multi-valued dependency occurs when one attribute determines another, and both attributes are independent of all other attributes in the table.

Example: Consider a table where (StudentID, Language, Hobby) are attributes. If a student can have multiple hobbies and languages, a multi-valued dependency exists. To resolve this, split the table into separate tables for Languages and Hobbies.

- Problem: A Restaurant table listing Cuisines and Delivery\_Areas. Cuisines and Areas are not related to each other but both relate to the Restaurant.
- SQL Example:

//4NF: Separate the independent multi-valued facts

```
CREATE TABLE Restaurant_Cuisines (RestID INT, Cuisine VARCHAR(50));
```

```
CREATE TABLE Restaurant_Areas (RestID INT, Area VARCHAR(50));
```

## 6. Fifth Normal Form (5NF): Eliminating Join Dependency

5NF is achieved when a table is in 4NF and all join dependencies are removed. This form ensures that every table is fully decomposed into smaller tables that are logically connected without losing information.

Example: If a table contains (StudentID, Course, Instructor) and there is a dependency where all combinations of these columns are needed for a specific relationship, you would split them into smaller tables to remove redundancy.

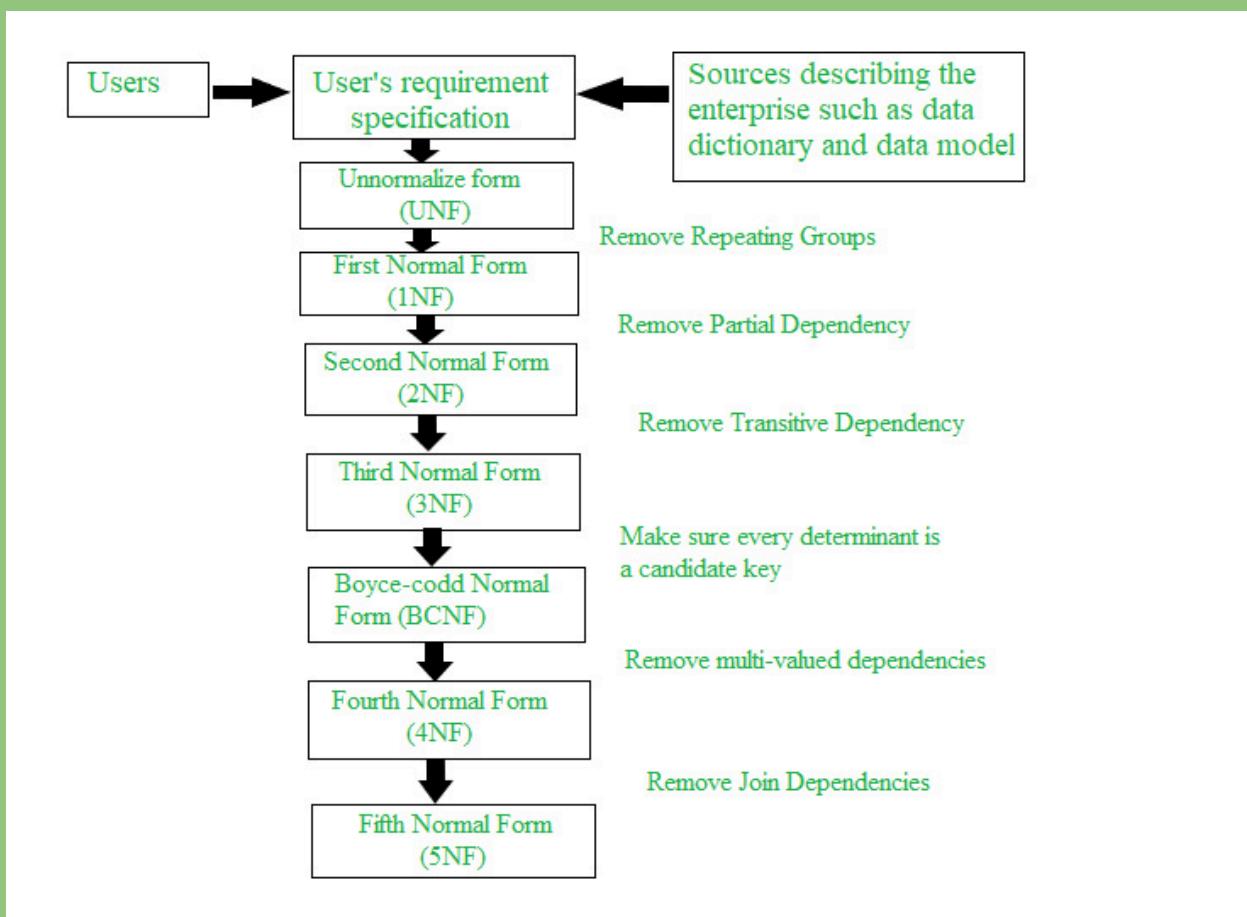
- Problem: A relationship between Brand, Product, and Salesman where all three must be linked to represent a specific business rule.
- SQL Example:

//5NF: Breaking down a 3-way relationship into 3 separate join tables

```
CREATE TABLE Brand_Product (BrandID INT, ProductID INT);
```

```
CREATE TABLE Product_Salesman (ProductID INT, SalesmanID INT);
```

```
CREATE TABLE Brand_Salesman (BrandID INT, SalesmanID INT);
```



## Common Challenges of Over-Normalization

While normalization is a powerful tool for optimizing databases, it's important not to over-normalize your data. Excessive normalization can lead to:

**Complex Queries:** Too many tables may result in multiple joins, making queries slow and difficult to manage.

**Performance Overhead:** Additional processing required for joins in overly normalized databases may hurt performance, especially in large-scale systems.

In many cases, denormalization (combining tables to reduce the need for complex joins) is used for performance optimization in specific applications, such as reporting systems.

## When to Use Normalization and Denormalization

Normalization is best suited for transactional systems where data integrity is paramount, such as banking systems and enterprise applications.

**Denormalization** is ideal for read-heavy applications like data warehousing and reporting systems where performance and query speed are more critical than data integrity.

## Applications of Normal Forms in DBMS

- Ensures Data Consistency: Prevents data anomalies by ensuring each piece of data is stored in one place, reducing inconsistencies.
- Reduces Data Redundancy: Minimizes repetitive data, saving storage space and avoiding errors in data updates or deletions.
- Improves Query Performance: Simplifies queries by breaking large tables into smaller, more manageable ones, leading to faster data retrieval.
- Enhances Data Integrity: Ensures that data is accurate and reliable by adhering to defined relationships and constraints between tables.
- Easier Database Maintenance: Simplifies updates, deletions, and modifications by ensuring that changes only need to be made in one place, reducing the risk of errors.
- Facilitates Scalability: Makes it easier to modify, expand, or scale the database structure as business requirements grow.
- Supports Better Data Modeling: Helps in designing databases that are logically structured, with clear relationships between tables, making it easier to understand and manage.
- Reduces Update Anomalies: Prevents issues like insertion, deletion, or modification anomalies that can arise from redundant data.
- Improves Data Integrity and Security: By reducing unnecessary data duplication, normal forms help ensure sensitive information is securely and correctly maintained.
- Optimizes Storage Efficiency: By organizing data into smaller tables, storage is used more efficiently, reducing the overhead for large databases

## **Example of Normalization**

Imagine a table storing student projects: [Student\_ID, Project\_ID, Project\_Name, Student\_Name].

Unnormalized: If a student works on three projects, their name is repeated three times.

Normalized (2NF/3NF): Break this into two tables: [Student\_ID, Student\_Name] and [Project\_ID, Project\_Name, Student\_ID]. This removes the redundancy.

## **Advantages and Disadvantages**

Pros: Better data integrity, higher consistency, and an organized structure.

Cons: Increased query complexity due to multiple joins and potential performance overhead.